SafePath Analytics:

Privacy-Preserving Auditing of Algorithmic Feedback Loops

Lead Architect: Yurun Liu
Document Version: 1.0 (Open-Source Architecture Draft)

Executive Summary & Methodological Anchor

Executive Summary

This paper proposes a decentralized, privacy-preserving analytics framework for auditing algorithmic feedback loops in educational environments. By combining local differential privacy with advanced causal inference, the **SafePath Analytics** toolkit bridges the massive information asymmetry between proprietary social media platforms and the public sector. While digital platforms utilize complex recommender systems aggressively optimized for engagement, public institutions (schools, youth NGOs, policymakers) lack the technical infrastructure to independently audit long-term well-being outcomes. SafePath disrupts this market failure by providing schools with a regulatory-grade mechanism to scientifically measure and mitigate digital harm without relying on platform self-reporting.

The Problem Space: Why Traditional Firewalls Fail

Current digital safety measures in educational environments are fundamentally reactive and inadequate due to three systemic bottlenecks:

- **The Black Box Effect:** Public institutions cannot link content exposure to downstream behavioral outcomes without access to proprietary platform data.
- **Lack of Causal Rigor:** Existing school monitoring tools rely on crude correlations (e.g., DNS blocking) rather than causal inference, failing to identify algorithmically induced rabbit holes.
- **The Privacy Bottleneck:** Analyzing student digital behavior traditionally risks violating strict privacy laws (e.g., FERPA/COPPA), preventing meaningful data collection at the local level.

Methodological Anchor: The Analytics Weapon

Unlike traditional web-filters, SafePath deploys advanced data science directly at the local edge. It bridges the technical gap by utilizing:

- **Propensity Score Matching (PSM) & Difference-in-Differences (DiD):** To neutralize confounding biases in observational data and isolate the true *causal* impact of algorithmic exposure.
- **Small-Sample Causal Estimation:** Educational pilots often face high-uncertainty

and constrained sample sizes. The framework utilizes quasi-experimental designs and robust small-sample causal estimators to derive high-confidence ($p < 0.05$) intervention signals without requiring massive, platform-scale datasets.

- **ϵ-Differential Privacy (ϵ-DP):** To mathematically guarantee FERPA compliance by irreversibly sanitizing telemetry before any causal modeling occurs.

Core Architecture & The FERPA Privacy Shield

Zero-Trust Architecture Flow

The SafePath architecture is fundamentally designed around a "Zero-Trust Data Transfer" principle. The pipeline consists of three isolated layers to ensure absolute data sovereignty for public schools.

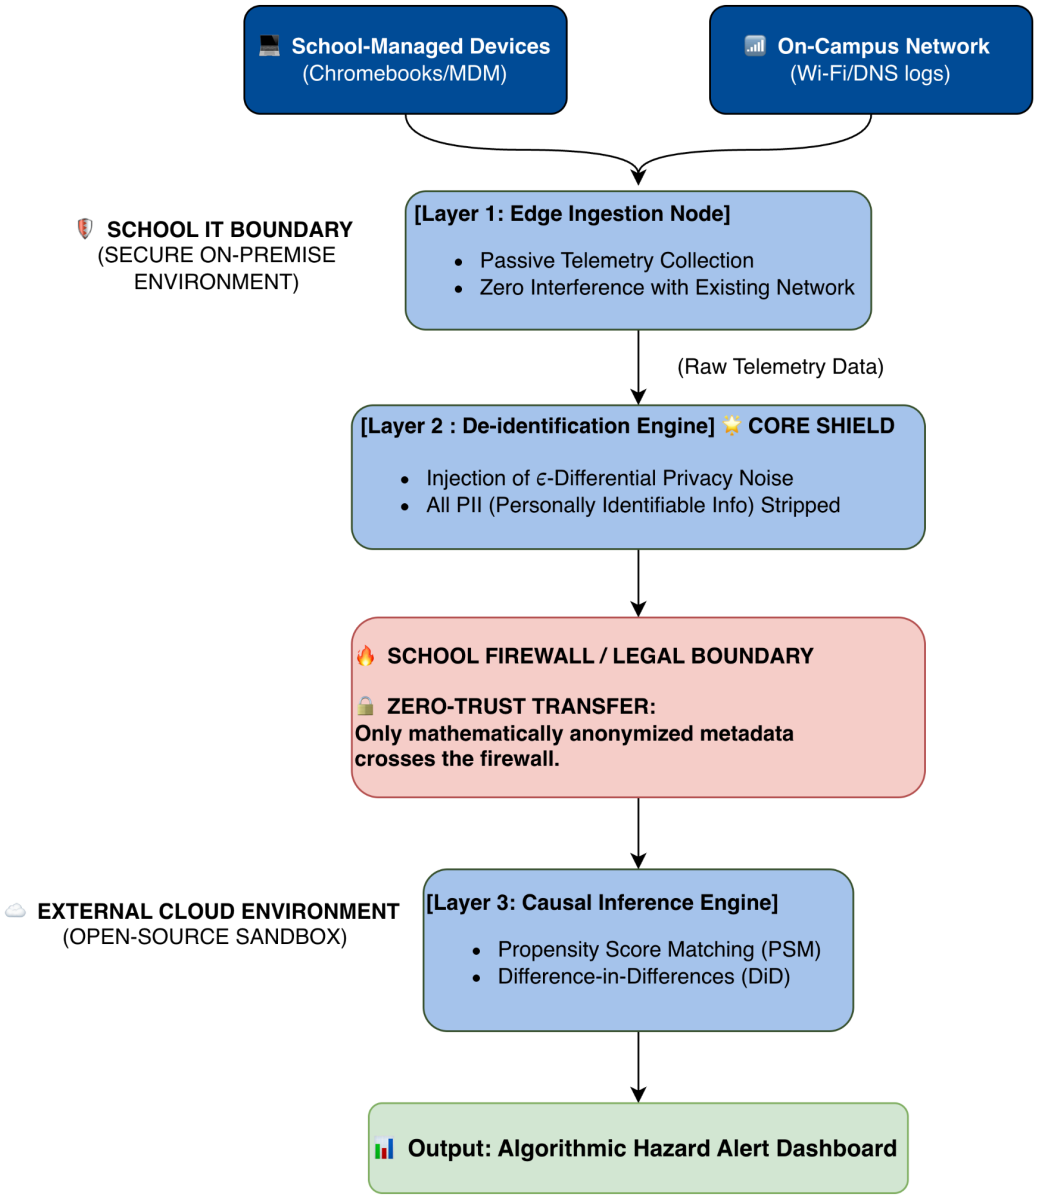

Layer 1: The Edge Ingestion Node (On-Premise)

- **Data Source:** Strictly limits its scope to School-Managed Infrastructure (MDM Telemetry & On-Campus Network Logs). We analyze the subset of digital exposure that occurs within the educational environment to extrapolate broader algorithmic feedback loops.
- **Function:** Deployed locally behind the school district's firewall. It intercepts raw web logs and API telemetry but *never* transmits them externally. It is designed as a lightweight, plug-and-play module requiring zero new hardware procurement.

Layer 2: De-identification Engine (The FERPA Shield)

- **Function:** Raw data points (IP addresses, student IDs) are stripped and irreversibly hashed. We apply **Local Differential Privacy (LDP)** algorithms at this edge layer, injecting calibrated Laplace noise into the data.
- **Compliance:** We mathematically guarantee that the inclusion or exclusion of any single student's digital behavior does not significantly alter the aggregated output. **No identifiable student data ever leaves the school's physical premises.** This guarantees 100% FERPA and COPPA compliance, ensuring the school assumes zero legal liability for individual data breaches.

Layer 3: Causal Auditing Sandbox (Cloud Environment)

- **Function:** Only fully anonymized, aggregated metadata (e.g., Cohort A experienced a 40% spike in late-night short-video traffic) reaches the open-source causal inference engine to compute Algorithmic Harm Indices.

Core Use Case, Scalability, & National Impact

Core Use Case in Action: The "Late-Night Rabbit Hole" Audit

To illustrate the pipeline's capability, consider a public school district seeking to understand the impact of infinite-scroll features on student sleep proxies and morning attention deficits.

- **The Input:** The local edge node securely ingests DP-sanitized network telemetry, identifying a structural shift in algorithmic delivery (e.g., a transition from educational streaming to extreme high-arousal short videos after 10:00 PM on school-issued Chromebooks).
- **The Analytics (Causal Engine):** The toolkit isolates this sequence. Using **Inverse Probability Weighting (IPW)** and **Difference-in-Differences (DiD)** estimation, it compares the downstream engagement duration against a synthetic control group not exposed to the loop, neutralizing baseline behavioral differences.
- **The Output:** Instead of simply blocking a URL, the district receives a statistically rigorous $p < 0.01$ Algorithmic Hazard Alert. This provides administrators with

empirical, legally defensible evidence to justify local firewall throttling of specific proprietary recommendation loops.

Implementation Roadmap & National Scalability

The true power of this open-source initiative lies in its scalability to address a national crisis.

- **Phase 1 (MVP):** Deployment of the local telemetry ingestion node and visualization dashboard for early-risk indicator tracking within pilot public schools.
- **Phase 2 (Content Benchmarking):** Integration of multi-modal LLM classifiers to evaluate the semantic safety of recommendation feedback loops without human-in-the-loop privacy breaches.
- **Phase 3 (Federated Learning Grid):** Enabling cross-district auditing. Multiple school districts (e.g., SFUSD and LAUSD) can collaboratively train Algorithmic Hazard Classifiers without ever sharing underlying student data, building a decentralized, nationwide defense grid against predatory algorithms.